

SNEWS Operational Mode 3.1 DRAFT

Effective [whenever backup setup is complete]

This document provides specifications for the instance of the class defined in the SNEWS Operational Mode Template document.

1 Participating Experiments

The participating experiments for this mode are:

- Super-K,
- LVD,
- AMANDA/IceCube,

Note that SNO members remain active members of the SNEWS working group.

2 Privacy Agreement

The SNEWS inter-experiment privacy agreement can be found at <http://snews.bnl.gov/wg/privacy/>.

3 Client Computer Management

Criteria of the privacy agreement (section 2) regarding access to SNEWS-related accounts on the client computer should be strictly observed. In addition, the following are recommended:

- Outside the main domain of the client computer, TCP/IP access to the client machine should be restricted.
- Login to the client computer should be permitted only to SNEWS subgroup members.
- The TCP/IP address of the client computer should be known only to SNEWS subgroup members.

4 Server Computer Management

There is one coincidence server machine located at Brookhaven National Laboratory, and the BNL SNEWS sysadmin is Brett Viren. A backup server machine is also in place at BNL for the case of hardware failure at BNL. There is also a backup server machine located at INFN, Sezione di Bologna, and the INFNBO sysadmin is Andrea Paolucci. The BNL server is the main server, and to the extent possible will operate at all times. The INFNBO server will be kept running in parallel with the BNL server, but will normally have alert output disabled; its alert output will be enabled only in the case of BNL network communications failure. Both servers will normally accept datagrams from all client machines.

The details of the server management, which address all security and maintenance issues for both machines, can be found in the Memorandum of Understanding at

http://snews.bnl.gov/wg/server_management/.

5 Client-Server Communications

- Each participating experiment may generate and send to a server different types of alarm datagrams. Each datagram will normally be sent to both the primary BNL server and the INFNBO backup server. The alarm datagrams include a packet type, and a level flag. The packet type can be PING, ALARM, or RETRACTION. The level flag can be TEST, GOOD, POSSIBLE, RETRACTED or CONFIRMED. Datagrams having packet field values which do not belong to any of these categories are discarded by the server.

Packet Types

- PING: Ping packets are used for test purposes only and cause nothing more than a message printed to the coincidence server log.
- ALARM: Alarm packets contain information about individual experiment alarms; what the server does with them depends on the level flag.
- RETRACTION: Retraction packets contain information about previously sent alarms to be retracted from the server's alarm queues.

Level Flags

- TEST: This flag indicates a datagram packet intended for test use as well as for any high-rate test mode. The alarm level is set to 0.
- POSSIBLE : This flag indicates an alarm generated during scheduled operations (i.e. maintenance, calibration, tests, etc.) or other known anomalous conditions. It is up to each experiment to set this flag inside the packet when appropriate. The alarm level is set to 1.

- GOOD: This flag indicates an alarm generated during normal detection conditions. The alarm level is set to *2*.
 - RETRACTED: This flag is set for retraction packets (note that this information is redundant – all packets of RETRACTION type will be retracted regardless of level flag). The alarm level is set to *-1*.
 - CONFIRMED: This flag indicates an alarm that has been confirmed as good. The alarm level is set to *3*.
- The details of the SNEWS datagram structure should not be propagated beyond the SNEWS subgroup.
 - The client datagrams employ TCP protocol and are encrypted using OpenSSL. The server process listens for connections on a specified port, and when a client initiates a connection, it employs several layers of checks to validate the origin of the datagram:
 1. The server process employs a `tcp_wrappers` `hosts_access` call to check the IP of the client machine against the lists in `hosts.allow/deny`. These lists allow only the IP addresses of the client machines of the involved experiments to submit packets.
 2. The client and server exchange certificates which have been verified by the SNEWS Certificate Authority, and reject connections if the check fails.
 3. There is an additional check of specific client machine IP addresses against a list of allowed addresses before data is exchanged.

6 SNEWS Shift

SNEWS shifts will be coordinated by Kate Scholberg, and will be run on a week by week basis. SNEWS shifters must be subgroup members. The shifter must have cellular communications capability, and have quick access to a computer during the 24 hour period of his or her shift.

6.1 Shift Alert Responsibilities

Upon receipt of any coincidence message, either SILVER or GOLD, the shifter must immediately log on to the server computer to check that SNEWS code is running correctly, according to logs. If everything is in order, he or she must follow the alert procedures outlined in sections 8.3-8.4. If something seems to be wrong, the shiftworker must contact the relevant people.

6.2 Shift Maintenance and Monitoring Responsibilities

The shifter has in addition some daily monitoring responsibilities. The online shift checklist is available at

<https://snews.bnl.gov/wg/shift/>. Twice per 24 hour period (morning and evening), the shifter must check the following:

- The gcserver process is running normally on both the BNL server and INFNBO backup server machines.
- The server log output seems normal on both the BNL server and INFNBO backup server machines.
- Ping connections from the shifter’s client machine to both the BNL server and INFNBO backup server machines work.
- The shifter’s cellular communications capability is active, batteries charged, etc.

The shifter must make an entry in the online logbook and record any anomalous conditions, and, as appropriate, notify the rest of the subgroup of any problems. If any particular experiment has an unusually high rate of alarms, the shifter should notify that experiment’s subgroup member.

If the shifter finds that the network connection to the primary server is unavailable, he or she must notify the relevant sysadmins, and initiate the failover procedure to the backup server.

7 Coincidence Definition

The general coincidence definition implemented in the coincidence code may generate either of two types of alert: GOLD or SILVER. An INDIVIDUAL alert may also be generated from a single experiment.

A GOLD alert is generated if *all* of following conditions (1 through 4) are met:

1. There is a 2 or more -fold coincidence within 10 seconds, involving at least two different experiments. (The time window refers to the maximum separation of any of the alarms in the coincidence.)
2. At least two of the experiments involved are at different laboratories. This condition is automatically satisfied for the current operational mode.
3. Two or more of the alarms in the coincidence are flagged as GOOD. It is the responsibility of each participant experiment to flag the alarm sent to the SNEWS server(s) appropriately. The specific criteria for a GOOD/POSSIBLE/TEST/CONFIRMED alarms are locally defined by each experiment according to the guidelines in section 5.
4. For at least two of the experiments involved in the coincidence, the rate of GOOD alarms for several past time intervals $\{T_i\} = \{10 \text{ minutes}, 1 \text{ hour}, 10 \text{ hours}, 1 \text{ day}, 3 \text{ days}, 1 \text{ week}, 1 \text{ month}\}$ preceding the first alarm of the coincidence candidate, must be consistent with the $\lambda_{\text{max}} = 1/7$ days

requirement.¹ We define the precise condition as follows: if an experiment sent $\{n_i\}$ alarms in each of the last intervals $\{T_i\}$ before the first event of the coincidence, then the Poisson probabilities \mathcal{P}_i for n_i or more alarms in T_i ,

$$\mathcal{P}_i = \sum_{n=n_i}^{\infty} (\lambda_{\max} T_i)^n e^{-\lambda_{\max} T_i} / n!,$$

for each interval T_i , must each be greater than $\mathcal{P}_{thr} = 0.5\%$. This corresponds to the condition that each $\{n_i\}$ must be less than $\{1, 2, 2, 3, 4, 5, 11\}$ for the preceding intervals $\{T_i\}$ for an alarm to be GOLD.

When the first criterion is satisfied, but at least one of the other criteria is not satisfied, the generated alert is flagged as SILVER. In this case the alert has to be checked by the individual experiment collaborations before any public announcement. No alert will be sent to the community by SNEWS until (and if) there is an upgrade to GOLD.

A GOLD alert may also be generated if condition 1 is satisfied and at least one alarm in the coincidence is CONFIRMED, regardless of whether the other conditions are satisfied. This case allows an override of past high rate history demotion, or other dubious conditions, for a human-checked alert.

An INDIVIDUAL alert is generated if a single experiment sends a CONFIRMED alarm.

8 Alert Procedure

8.1 The Alert Message

For both SILVER or GOLD coincidence cases, a message containing the following information will be automatically sent by the server to the SNEWS subgroup members:

- all detectors involved in the coincidence,
- the types of alarms (GOOD, POSSIBLE, CONFIRMED) for each experiment involved in the coincidence,
- UTC time of each alarm in the coincidence, and
- if available, information about the burst: duration (seconds), number of signal events, right ascension and declination (degrees) of the inferred supernova direction, and error on the direction (degrees), for the alarms in the coincidence. If this information is not available, default values of zero are given, and 360 degrees for the direction error. Note that this extra information need not be sent even if available; it is sent at the discretion of the experiment's collaboration.

¹These intervals represent real time, not live time, since full live time information will not be available to the coincidence server.

The information may also be posted to a restricted SNEWS subgroup page for SILVER, and a public page for GOLD.

For an INDIVIDUAL alert, the same information as above is sent for the single CONFIRMED level experiment.

An example of the alert message follows, for which extra information is available for one out of the three experiments in the coincidence; the other experiments have default values.

Date: Wed, 17 Mar 2004 09:46:24 -0500
From: SN Alert <snalert@bnlboom1.bnl.gov>
To: an.observer@mydomain.com
Subject: Gold Coincidence Alert

*** SNEWS ALERT ***

Coincidence rating: GOLD
Alarms in the coincidence:
Experiment: 1 Super-K
Level: GOOD
Time: Mar 16 2004 16:04:04.55000000
Duration: 0.00
No. of signal events: 0.00
Right Ascension: 0.00
Declination: 0.00
Error: 360.00

Experiment: 3 SNO
Level: GOOD
Time: Mar 16 2004 16:04:05.02300000
Duration: 6.00
No. of signal events: 400.00
Right Ascension: 23.00
Declination: -82.00
Error: 10.00

Experiment: 5 LVD
Level: GOOD
Time: Mar 16 2004 16:04:05.10300000
Duration: 0.00
No. of signal events: 0.00
Right Ascension: 0.00
Declination: 0.00
Error: 360.00

For information, see web page <http://snews.bnl.gov/>

Note that the server process searches the past day's worth of its alarms in memory.

8.2 GOLD Procedure

For a GOLD alert, the alert message is sent directly to the `snews-alert` mailing list, which includes addresses of all who have signed up to received it. This includes astronomers, and the Sky & Telescope AstroAlert list.

8.2.1 Experiment Gold Procedures

- Super-K: see Appendix for contact list
- LVD: see Appendix for contact list
- AMANDA/IceCube: see Appendix for contact list

Demoting from GOLD:

Although we hope to avoid ever being in the situation where retraction of a GOLD alert is necessary, any experiment may reflag from GOOD (or POSSIBLE) to RETRACTED its own alarm after data checking. The server will then automatically reevaluate and reissue the alert based on alarms in the past day of its memory: the result may be still GOLD, demotion to SILVER, or no alert at all. For the latter case, the SNEWS subgroup is notified, and a RETRACTED alert will be issued to the same mailing list as for GOLD and posted on the public web page.

8.3 SILVER Procedure

For a SILVER alert, no automated alert message to the community is generated by the coincidence server. The procedure indicated by each experiment should be followed by the SNEWS shiftworker.

8.3.1 Experiment Silver Procedures

- Super-K: see Appendix for contact list
- LVD: see Appendix for contact list
- AMANDA/IceCube: see Appendix for contact list

Upgrading from SILVER to GOLD

A SILVER alert may be upgraded to GOLD if the requirements are confirmed in a second step. Any individual experiment is expected to resend its alarm datagram if the flag can be changed from POSSIBLE to CONFIRMED. All subgroup members must be then informed by e-mail. If the GOLD conditions

are consequently fulfilled after the new alarm dispatch, the GOLD procedure will automatically operate. Note that the server process searches for coincidences in the past day's worth of alarms in memory.

8.4 INDIVIDUAL Procedure

The INDIVIDUAL procedure is identical to the GOLD one.

8.5 SNEWS Data Exchange

All server-client information exchanged during a supposed SN alert will be saved in the server log file. This information is available to subgroup members at any time, and later to the experiment collaborations by agreement of the SNEWS Advisory Board.

8.6 Summary

Figure 1 summarizes the sequence of events and GOLD vs. SILVER decisions.

9 Communication to the Scientific Community

9.1 Alert Mailing List Management

To allow the confirmation of a SNEWS alert as really coming from SNEWS, any alerts will be public key signed using the SNEWS key. This key has the ID# 0x68DF93F7, and is available on the network of public PGP key servers such as <http://pgp.mit.edu/>

People receiving any alert should verify that this signature was used to sign the alert message, to satisfy themselves that the alert is really coming from the SNEWS collaboration and is not being maliciously or accidentally spoofed. Exactly how to do this is a function of the individual user's email client. Since many people do not use such features on a day to day basis, anyone wishing to be able to verify SNEWS alerts should set up their client in advance. Instructions for doing so can be found on the SNEWS web site for many popular email clients.

Email to the list of alert subscribers will be limited to real alerts and quarterly test messages. These test messages will be clearly labeled as a mailing list test, to remind the subscribers that they are subscribed and to ensure mailing list functionality.

People wishing to subscribe or unsubscribe from the alert mailing list can do so automatically using the GNU mailman software linked to from the SNEWS website. The subscriber list is held in complete confidentiality, only available to the SNEWS core group and not used for any purpose other than alerts and periodic tests.

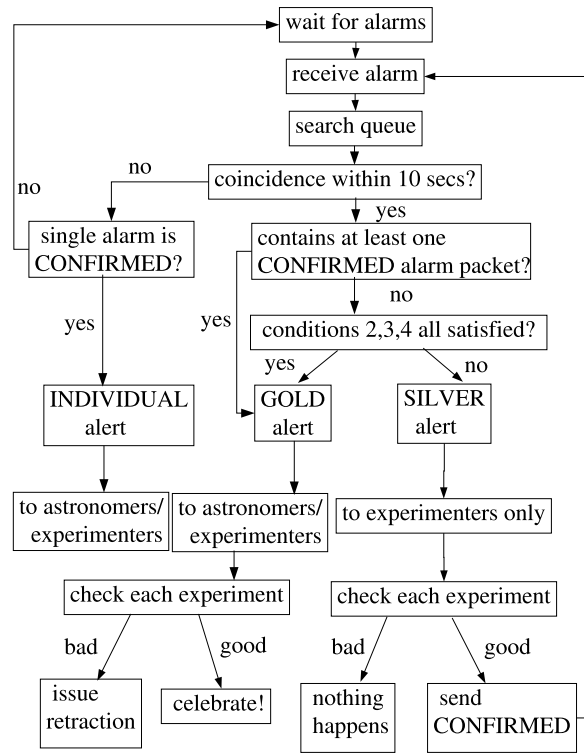


Figure 1: Flowchart summarizing the sequence of events and decisions that determine whether an alert is GOLD, SILVER (or INDIVIDUAL).

For the GOLD and INDIVIDUAL cases, the same information and updates will go to the public SNEWS URL. For the SILVER case, the message goes out only to the SILVER list.

9.2 Individual Experiment Alerts

The INDIVIDUAL alert mode is intended to allow individual experiments to make use of SNEWS alert infrastructure; it is in no way required. There is no restriction on individual experiments making any independent announcement based on individual observation in the case of absence of a SNEWS alert, SILVER or GOLD, or preceding or following any SNEWS alert message. Any individual experiment may publicly announce a supposed supernova signal following a dispatched SILVER alert which has not yet been upgraded to GOLD, or a GOLD alert. In these cases the information that a previous alert from the SNEWS server has been received should be cited.

10 Changelog

Version 3.1

- SNO is now off; gold alert criteria adjusted for 2/3 coincidence: mean alarm rate is changed from 1/week to 1/10 days; Poisson probability threshold for automatic rate suppression is changed from 0.4% to 0.5%.
- Describe new backup server, and criteria for failover.
- Add backup server check responsibilities to shift section.
- Update links to privacy and server management web pages.

Version 3.0

- Allow for INDIVIDUAL alert for a single CONFIRMED level alarm. Change flowchart accordingly.
- Send GOLD coincidence if a single CONFIRMED is included, regardless of other condition (previously at least one GOOD and one CONFIRMED was required).
- Change nomenclature: OVERRIDE \rightarrow CONFIRMED.

Version 2.1

- Specify additional optional allowed fields in the alarm (duration, number of signal events, direction.)

Version 2.0

- Add AMANDA/IceCube.

- Update web links.
- Small quantitative changes to criteria to accommodate the fact that it's now a 2/4 coincidence: mean alarm rate is changed from 1/week to 1/10 days; Poisson probability threshold for automatic rate suppression is changed from 0.5% to 0.4%.

Version 1.0

- First version, for SK/SNO/LVD.